

PART A

1) Imported database “largeRelationsInsertFile.sql”

2)

```

MariaDB [University]> SET profiling = 1;
Query OK, 0 rows affected (0.00 sec)

MariaDB [University]> select * from student where name='wood';
+-----+-----+-----+-----+
| ID    | name | dept_name | tot_cred |
+-----+-----+-----+-----+
| 33791 | Wood | Civil Eng. | 92       |
| 39876 | Wood | Accounting | 14       |
| 62054 | Wood | Mech. Eng. | 13       |
| 96085 | Wood | Accounting | 70       |
+-----+-----+-----+-----+
4 rows in set (0.00 sec)

MariaDB [University]> SHOW PROFILES;
+-----+-----+-----+-----+
| Query_ID | Duration | Query |
+-----+-----+-----+-----+
| 1        | 0.00006858 | select * from student where name='wood' |
+-----+-----+-----+-----+
1 row in set (0.00 sec)

MariaDB [University]> SHOW PROFILE FOR QUERY 1;
+-----+-----+-----+-----+
| Status | Duration |
+-----+-----+-----+-----+
| starting | 0.000022 |
| Waiting for query cache lock | 0.000004 |
| init | 0.000003 |
| checking query cache for query | 0.000007 |
| checking privileges on cached | 0.000004 |
| checking permissions | 0.000008 |
| sending cached result to clien | 0.000013 |
| updating status | 0.000005 |
| cleaning up | 0.000003 |
+-----+-----+-----+-----+
9 rows in set (0.00 sec)

MariaDB [University]>

```

Thus the bottleneck is for starting and then comes the sending cached results

```

MariaDB [University]> SHOW PROCESSLIST;
+-----+-----+-----+-----+-----+-----+-----+-----+
| Id | User | Host | db | Command | Time | State | Info | Progress |
+-----+-----+-----+-----+-----+-----+-----+-----+
| 2 | system user | | NULL | Daemon | NULL | InnoDB purge worker | NULL | 0.000 |
| 1 | system user | | NULL | Daemon | NULL | InnoDB purge coordinator | NULL | 0.000 |
| 3 | system user | | NULL | Daemon | NULL | InnoDB purge worker | NULL | 0.000 |
| 4 | system user | | NULL | Daemon | NULL | InnoDB purge worker | NULL | 0.000 |
| 5 | system user | | NULL | Daemon | NULL | InnoDB shutdown handler | NULL | 0.000 |
| 11 | root | localhost | University | Query | 0 | init | SHOW PROCESSLIST | 0.000 |
+-----+-----+-----+-----+-----+-----+-----+-----+
6 rows in set (0.00 sec)

```

```

MariaDB [University]> show ENGINES\G;
+-----+-----+
| Engine | Support | Comment | Transactions | XA | Savepoints |
+-----+-----+
1. row +-----+
| CSV | YES | CSV storage engine | NO | NO | NO |
+-----+-----+
2. row +-----+
| MRG_MyISAM | YES | Collection of identical MyISAM tables | NO | NO | NO |
+-----+-----+
3. row +-----+
| SEQUENCE | YES | Generated tables filled with sequential values | YES | NO | YES |
+-----+-----+
4. row +-----+
| MyISAM | YES | MyISAM storage engine | NO | NO | NO |
+-----+-----+
5. row +-----+
| MEMORY | YES | Hash based, stored in memory, useful for temporary tables | NO | NO | NO |
+-----+-----+
6. row +-----+
| PERFORMANCE_SCHEMA | YES | Performance Schema | NO | NO | NO |
+-----+-----+
7. row +-----+
| Aria | YES | Crash-safe tables with MyISAM heritage | NO | NO | NO |
+-----+-----+
8. row +-----+
| InnoDB | DEFAULT | Supports transactions, row-level locking, foreign keys and encryption for tables | YES | YES | YES |
+-----+-----+
8 rows in set (0.00 sec)

```

Default storage machine is InnoDB

Engine MEMORY is Hash based hence support Hash Index

```

MariaDB [University]> CREATE TABLE `takes_hash` ENGINE=MEMORY AS (SELECT * FROM takes);
Query OK, 30000 rows affected (0.08 sec)
Records: 30000 Duplicates: 0 Warnings: 0

MariaDB [University]> describe takes_hash
-> ;
+-----+-----+-----+-----+-----+-----+
| Field | Type | Null | Key | Default | Extra |
+-----+-----+-----+-----+-----+-----+
| ID | varchar(5) | NO | | | |
| course_id | varchar(8) | NO | | | |
| sec_id | varchar(8) | NO | | | |
| semester | varchar(6) | NO | | | |
| year | decimal(4,0) | NO | | 0 | |
| grade | varchar(2) | YES | | NULL | |
+-----+-----+-----+-----+-----+-----+
6 rows in set (0.00 sec)

MariaDB [University]> CREATE INDEX grade_index ON takes_hash(grade) USING HASH;
Query OK, 30000 rows affected (0.05 sec)
Records: 30000 Duplicates: 0 Warnings: 0

```

```
MariaDB [University]> SHOW INDEX FROM takes_hash;
```

Table	Non_unique	Key_name	Seq_in_index	Column_name	Collation	Cardinality	Sub_part	Packed	Null	Index_type	Comment	Index_comment
takes_hash	1	grade_index	1	grade	NULL	9	NULL	NULL	YES	HASH		

```
1 row in set (0.00 sec)
```

```
MariaDB [University]> SHOW INDEXES FROM takes_hash;
```

Table	Non_unique	Key_name	Seq_in_index	Column_name	Collation	Cardinality	Sub_part	Packed	Null	Index_type	Comment	Index_comment
takes_hash	1	grade_index	1	grade	NULL	9	NULL	NULL	YES	HASH		

```
1 row in set (0.00 sec)
```

```
MariaDB [University]>
```

```
MariaDB [University]> CREATE INDEX h_gr ON takes(grade) USING HASH;
Query OK, 0 rows affected (1.36 sec)
Records: 0 Duplicates: 0 Warnings: 0
```

```
MariaDB [University]> SHOW INDEXES FROM takes;
```

Table	Non_unique	Key_name	Seq_in_index	Column_name	Collation	Cardinality	Sub_part	Packed	Null	Index_type	Comment	Index_comment
takes	0	PRIMARY	1	ID	A	3984	NULL	NULL		BTREE		
takes	0	PRIMARY	2	course_id	A	27892	NULL	NULL		BTREE		
takes	0	PRIMARY	3	sec_id	A	27892	NULL	NULL		BTREE		
takes	0	PRIMARY	4	semester	A	27892	NULL	NULL		BTREE		
takes	0	PRIMARY	5	year	A	27892	NULL	NULL		BTREE		
takes	1	course_id	1	course_id	A	162	NULL	NULL		BTREE		
takes	1	course_id	2	sec_id	A	178	NULL	NULL		BTREE		
takes	1	course_id	3	semester	A	178	NULL	NULL		BTREE		
takes	1	course_id	4	year	A	178	NULL	NULL		BTREE		
takes	1	h_gr	1	grade	A	2	NULL	NULL	YES	BTREE		

```
10 rows in set (0.00 sec)
```

```
MariaDB [University]> SHOW INDEX FROM takes;
```

Thus hash is not used instead it uses a Btree index type which is nullable

```
MariaDB [University]> SHOW PROFILES;
```

Query_ID	Duration	Query
10	0.05188911	CREATE INDEX grade_index ON takes_hash(grade) USING HASH
11	0.00004350	show INDEX on FROM takes_hash
12	0.00007812	show INDEXES on FROM takes_hash
13	0.00042339	SHOW INDEX FROM takes_hash
14	0.00008978	SHOW INDEXES FROM takes_hash
15	0.00036734	SHOW INDEX FROM takes_hash
16	0.00040729	SHOW INDEXES FROM takes_hash
17	0.00008329	CREATE INDEX 'h_gr' ON takes_hash(grade) USING HASH
18	0.00008364	CREATE INDEX 'h_gr' ON takes(grade) USING HASH
19	1.35311463	CREATE INDEX h_gr ON takes(grade) USING HASH
20	0.00034350	SHOW INDEXES FROM takes
21	0.00034375	SHOW INDEX FROM takes
22	0.00010532	SET profiling = 1
23	0.01455390	select * from takes where grade LIKE '%C%'
24	0.01400009	select * from takes_hash where grade LIKE '%C%'

```
15 rows in set (0.00 sec)
```

That is hash takes less time than B tree Indexing.

MariaDB [University]> SHOW PROFILE FOR QUERY 23;

Status	Duration
starting	0.000020
Waiting for query cache lock	0.000006
init	0.000003
checking query cache for query	0.000038
checking permissions	0.000006
Opening tables	0.000018
After opening tables	0.000005
System lock	0.000004
Table lock	0.000005
Waiting for query cache lock	0.000016
init	0.000023
optimizing	0.000010
statistics	0.000014
preparing	0.000018
executing	0.000003
Sending data	0.008017
Waiting for query cache lock	0.000002
Sending data	0.000293
Waiting for query cache lock	0.000002
Sending data	0.000285
Waiting for query cache lock	0.000001
Sending data	0.000302
Waiting for query cache lock	0.000018
Sending data	0.000286
Waiting for query cache lock	0.000001
Sending data	0.000285
Waiting for query cache lock	0.000001
Sending data	0.000301
Waiting for query cache lock	0.000002
Sending data	0.000258
Waiting for query cache lock	0.000001
Sending data	0.000346
Waiting for query cache lock	0.000002
Sending data	0.000270
Waiting for query cache lock	0.000002
Sending data	0.000276
Waiting for query cache lock	0.000002
Sending data	0.000263
Waiting for query cache lock	0.000001
Sending data	0.000265
Waiting for query cache lock	0.000001
Sending data	0.000264
Waiting for query cache lock	0.000001
Sending data	0.000267
Waiting for query cache lock	0.000002
Sending data	0.000266
Waiting for query cache lock	0.000002

```
MariaDB [University]> SHOW PROFILE FOR QUERY 24;
```

Status	Duration
starting	0.000027
Waiting for query cache lock	0.000006
init	0.000004
checking query cache for query	0.000048
checking permissions	0.000008
Opening tables	0.000022
After opening tables	0.000005
System lock	0.000004
Table lock	0.000007
Waiting for query cache lock	0.000014
init	0.000028
optimizing	0.000013
statistics	0.000018
preparing	0.000024
executing	0.000005
Sending data	0.000723
Waiting for query cache lock	0.000006
Sending data	0.001082
Waiting for query cache lock	0.000008
Sending data	0.000742
Waiting for query cache lock	0.000005
Sending data	0.000939
Waiting for query cache lock	0.000034
Sending data	0.000840
Waiting for query cache lock	0.000008
Sending data	0.000732
Waiting for query cache lock	0.000006
Sending data	0.000727
Waiting for query cache lock	0.000006
Sending data	0.000707
Waiting for query cache lock	0.000006
Sending data	0.000710
Waiting for query cache lock	0.000006

```
MariaDB [University]>
```

```
MariaDB [University]> CREATE UNIQUE INDEX grade_index ON takes(grade);
```

```
ERROR 1062 (23000): Duplicate entry 'B+' for key 'grade_index'
```

```
MariaDB [University]>
```

```
MariaDB [University]> ALTER TABLE takes ADD INDEX compound(ID, course_id);
Query OK, 0 rows affected (1.37 sec)
Records: 0  Duplicates: 0  Warnings: 0
```

```
MariaDB [University]> SHOW INDEXES FROM takes;
```

Table	Non_unique	Key_name	Seq_in_index	Column_name	Collation	Cardinality	Sub_part	Packed	Null	Index_type	Comment	Index_comment
takes	0	PRIMARY	1	ID	A	4341	NULL	NULL		BTREE		
takes	0	PRIMARY	2	course_id	A	30392	NULL	NULL		BTREE		
takes	0	PRIMARY	3	sec_id	A	30392	NULL	NULL		BTREE		
takes	0	PRIMARY	4	semester	A	30392	NULL	NULL		BTREE		
takes	0	PRIMARY	5	year	A	30392	NULL	NULL		BTREE		
takes	1	course_id	1	course_id	A	162	NULL	NULL		BTREE		
takes	1	course_id	2	sec_id	A	178	NULL	NULL		BTREE		
takes	1	course_id	3	semester	A	178	NULL	NULL		BTREE		
takes	1	course_id	4	year	A	178	NULL	NULL		BTREE		
takes	1	h_gr	1	grade	A	42	NULL	NULL	YES	BTREE		
takes	1	compound	1	ID	A	4341	NULL	NULL		BTREE		
takes	1	compound	2	course_id	A	30392	NULL	NULL		BTREE		

```
12 rows in set (0.00 sec)
```

```
MariaDB [University]> \
```

4)

```
MariaDB [University]> SHOW VARIABLES WHERE Variable_name LIKE "datadir";
```

Variable_name	Value
datadir	/var/lib/mysql/

```
1 row in set (0.00 sec)
```

```
MariaDB [University]> SHOW VARIABLES WHERE Variable_name LIKE "%engine%";
```

Variable_name	Value
default_storage_engine	InnoDB
default_tmp_storage_engine	
enforce_storage_engine	
storage_engine	InnoDB

```
4 rows in set (0.00 sec)
```

```
MariaDB [University]> SHOW VARIABLES WHERE Variable_name LIKE "%buffer_size%";
```

Variable_name	Value
aria_pagecache_buffer_size	134217728
aria_sort_buffer_size	268434432
bulk_insert_buffer_size	8388608
innodb_log_buffer_size	16777216
innodb_sort_buffer_size	1048576
join_buffer_size	262144
key_buffer_size	16777216
mrr_buffer_size	262144
mysam_sort_buffer_size	134216704
preload_buffer_size	32768
read_buffer_size	131072
read_rnd_buffer_size	262144
sort_buffer_size	2097152

```
13 rows in set (0.00 sec)
```

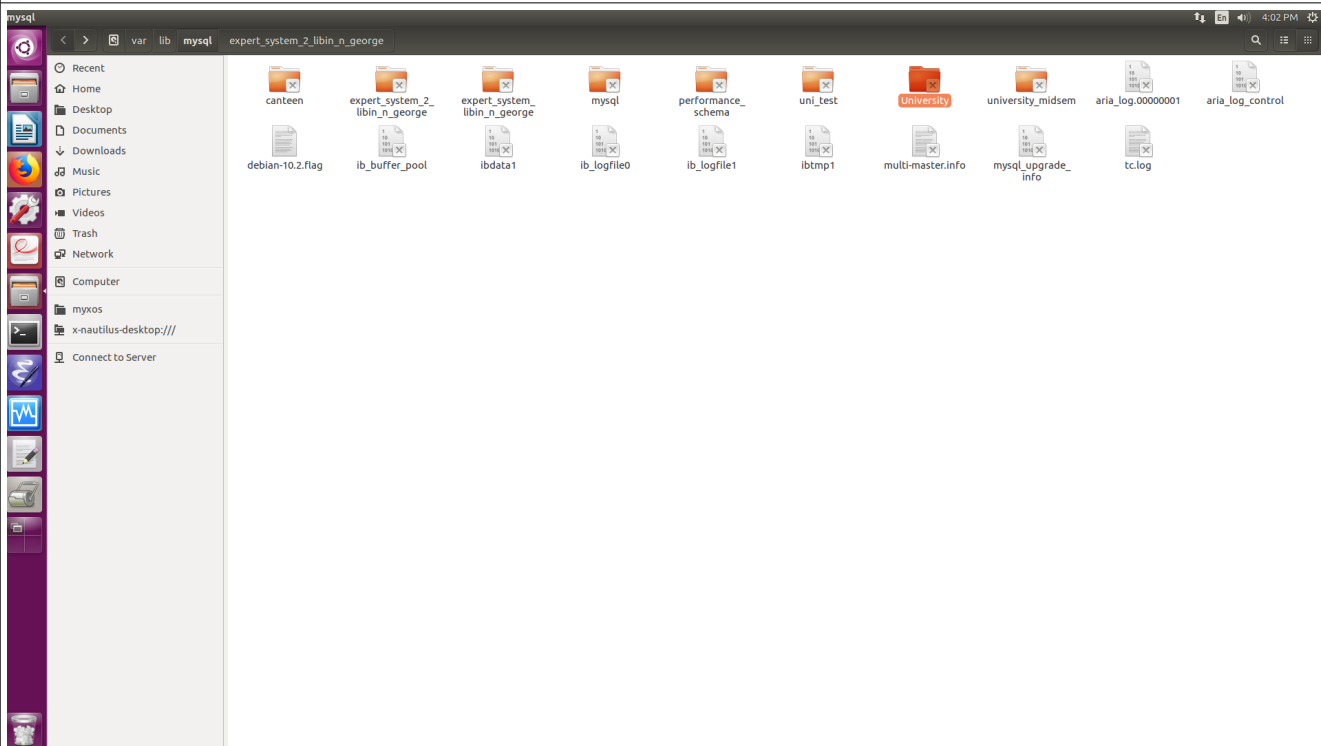
```

MariaDB [University]> set sort_buffer_size = 1097152;
Query OK, 0 rows affected (0.00 sec)

MariaDB [University]> SHOW VARIABLES WHERE Variable_name LIKE "%buffer_size%";
+-----+-----+
| Variable_name | Value |
+-----+-----+
| aria_pagecache_buffer_size | 134217728 |
| aria_sort_buffer_size | 268434432 |
| bulk_insert_buffer_size | 8388608 |
| innodb_log_buffer_size | 16777216 |
| innodb_sort_buffer_size | 1048576 |
| join_buffer_size | 262144 |
| key_buffer_size | 16777216 |
| mrr_buffer_size | 262144 |
| myisam_sort_buffer_size | 134216704 |
| preload_buffer_size | 32768 |
| read_buffer_size | 131072 |
| read_rnd_buffer_size | 262144 |
| sort_buffer_size | 1097152 |
+-----+-----+
13 rows in set (0.00 sec)

MariaDB [University]>

```



We can see each database as a inode/directory

PART B

```
> show dbs
admin  0.078GB
local  0.078GB
> use inventory;
switched to db inventory
>
```

```
> db.store.insert ([{item: "journal", qty: 25, size: { h: 14, w: 21, uom: "cm" }, status: "A"}, {item: "notebook", qty: 50, size: { h: 8.5, w: 11, uom: "in" }, status: "A"}, {item: "paper", qty: 100, size: { h: 8.5, w: 11, uom: "in" }, status: "D"}, {item: "planner", qty: 75, size: { h: 22.85, w: 30, uom: "cm" }, status: "D"}, {item: "postcard", qty: 45, size: { h: 10, w: 15.25, uom: "cm" }, status: "A"}])
BulkWriteResult({
  "writeErrors" : [ ],
  "writeConcernErrors" : [ ],
  "nInserted" : 5,
  "nUpserted" : 0,
  "nMatched" : 0,
  "nModified" : 0,
  "nRemoved" : 0,
  "upserted" : [ ]
})
> db.store.find().pretty()
{
  "_id" : ObjectId("5aa119ff22e2a024033c79c3"),
  "item" : "journal",
  "qty" : 25,
  "size" : {
    "h" : 14,
    "w" : 21,
    "uom" : "cm"
  },
  "status" : "A"
}
```

```
{
  "_id" : ObjectId("5aa11a7322e2a024033c79c4"),
  "item" : "journal",
  "qty" : 25,
  "size" : {
    "h" : 14,
    "w" : 21,
    "uom" : "cm"
  },
  "status" : "A"
}
{
  "_id" : ObjectId("5aa11a7322e2a024033c79c5"),
  "item" : "notebook",
  "qty" : 50,
  "size" : {
    "h" : 8.5,
    "w" : 11,
    "uom" : "in"
  },
  "status" : "A"
}
{
  "_id" : ObjectId("5aa11a7322e2a024033c79c6"),
  "item" : "paper",
  "qty" : 100,
  "size" : {
    "h" : 8.5,
    "w" : 11,
    "uom" : "in"
  },
  "status" : "D"
}
{
```



```

{
  "_id" : ObjectId("5aa11a7322e2a024033c79c7"),
  "item" : "planner",
  "qty" : 75,
  "size" : {
    "h" : 22.85,
    "w" : 30,
    "uom" : "cm"
  },
  "status" : "D"
}
{
  "_id" : ObjectId("5aa11a7322e2a024033c79c8"),
  "item" : "postcard",
  "qty" : 45,
  "size" : {
    "h" : 10,
    "w" : 15.25,
    "uom" : "cm"
  },
  "status" : "A"
}
>

```

```

> db.store.find({"status":"D"}).pretty()
{
  "_id" : ObjectId("5aa11a7322e2a024033c79c6"),
  "item" : "paper",
  "qty" : 100,
  "size" : {
    "h" : 8.5,
    "w" : 11,
    "uom" : "in"
  },
  "status" : "D"
}
{
  "_id" : ObjectId("5aa11a7322e2a024033c79c7"),
  "item" : "planner",
  "qty" : 75,
  "size" : {
    "h" : 22.85,
    "w" : 30,
    "uom" : "cm"
  },
  "status" : "D"
}
>

```

```
> db.store.find({$or: [{"status":"D"}, {"status":"A"} ]}).pretty()
{
  "_id" : ObjectId("5aa119ff22e2a024033c79c3"),
  "item" : "journal",
  "qty" : 25,
  "size" : {
    "h" : 14,
    "w" : 21,
    "uom" : "cm"
  },
  "status" : "A"
}
{
  "_id" : ObjectId("5aa11a7322e2a024033c79c4"),
  "item" : "journal",
  "qty" : 25,
  "size" : {
    "h" : 14,
    "w" : 21,
    "uom" : "cm"
  },
  "status" : "A"
}
{
  "_id" : ObjectId("5aa11a7322e2a024033c79c5"),
  "item" : "notebook",
  "qty" : 50,
  "size" : {
    "h" : 8.5,
    "w" : 11,
    "uom" : "in"
  },
  "status" : "A"
}
```

```

}
{
  "_id" : ObjectId("5aa11a7322e2a024033c79c5"),
  "item" : "notebook",
  "qty" : 50,
  "size" : {
    "h" : 8.5,
    "w" : 11,
    "uom" : "in"
  },
  "status" : "A"
}
{
  "_id" : ObjectId("5aa11a7322e2a024033c79c6"),
  "item" : "paper",
  "qty" : 100,
  "size" : {
    "h" : 8.5,
    "w" : 11,
    "uom" : "in"
  },
  "status" : "D"
}
{
  "_id" : ObjectId("5aa11a7322e2a024033c79c7"),
  "item" : "planner",
  "qty" : 75,
  "size" : {
    "h" : 22.85,
    "w" : 30,
    "uom" : "cm"
  },
  "status" : "D"
}

```

```

> db.store.find({$and: [{"status": "A"}], {$or: [{"qty": {$gt: 30}}, {"item": {$regex: /p.*//}}]}]).pretty()

```

```

{
  "_id" : ObjectId("5aa11a7322e2a024033c79c5"),
  "item" : "notebook",
  "qty" : 50,
  "size" : {
    "h" : 8.5,
    "w" : 11,
    "uom" : "in"
  },
  "status" : "A"
}
{
  "_id" : ObjectId("5aa11a7322e2a024033c79c8"),
  "item" : "postcard",
  "qty" : 45,
  "size" : {
    "h" : 10,
    "w" : 15.25,
    "uom" : "cm"
  },
  "status" : "A"
}
>

```

```

libin@Lenovo-Yoga-500-14IBD:/media/libin/study/DataBase-Lab$ python con.py
Number of records in store is 6
{'_id': ObjectId('5aa11a7322e2a024033c79c6'),
 'item': u'paper',
 'qty': 100.0,
 'size': {'h': 8.5, 'uom': u'in', 'w': 11.0},
 'status': u'D'}
item name =paper
{'_id': ObjectId('5aa11a7322e2a024033c79c7'),
 'item': u'planner',
 'qty': 75.0,
 'size': {'h': 22.85, 'uom': u'cm', 'w': 30.0},
 'status': u'D'}
item name =planner

```

```

from pymongo import MongoClient
import pprint
client = MongoClient()
client = MongoClient("mongodb://localhost")
db = client.inventory
coll = db.store
print "Number of records in store is {}".format(coll.count())
for i in coll.find({"status":"D"}):
    pprint.pprint(i)
    print "item name ={} ".format(i["item"])

```

```

from pymongo import MongoClient
import pprint
client = MongoClient()
client = MongoClient("mongodb://localhost")
db = client.inventory
coll = db.store
di = []
for i in coll.find({"status":"D"}):
    di.append({"Item":i["item"],"Quantity": i["qty"]})
for i in di:
    print "Item = {},Quantity= {}".format( i["Item"], i["Quantity"])

```

```

libin@Lenovo-Yoga-500-14IBD:/
Item = paper,Quantity= 100.0
Item = planner,Quantity= 75.0
libin@Lenovo-Yoga-500-14IBD:/

```